

Basic Importing

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import math as m
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

Importing Data

In [2]:

```
california_data = pd.read_csv('california_housing_data.csv')#main data
california_cities = pd.read_csv('cal_cities_lat_long.csv')#city locations
```

Cleaning the Data

In [209]:

```
california_data.describe()
```

Out[209]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	p
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1443.870553
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1133.870553
min	-124.350000	32.540000	1.000000	2.000000	1.000000	1.000000
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	781.000000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1161.000000
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1721.000000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	3568.000000

note 20640 total rows, only 5% can be taken out giving 19608 rows still have to be included at the end of cleaning

In [210]:

```
california_data.isnull().sum()
```

Out[210]:

```
longitude          0
latitude           0
housing_median_age  0
total_rooms         0
total_bedrooms     207
population          0
households          0
median_income       0
median_house_value  0
ocean_proximity    0
dtype: int64
```

The cell below drops rows with missing values

In [4]:

```
housing_raw=california_data.dropna()
housing_raw.isnull().sum()
```

Out[4]:

```
longitude          0
latitude           0
housing_median_age  0
total_rooms         0
total_bedrooms     0
population          0
households          0
median_income       0
median_house_value  0
ocean_proximity    0
dtype: int64
```

In [212]:

```
print(housing_raw.nunique())
```

```
longitude          844
latitude           861
housing_median_age   52
total_rooms        5911
total_bedrooms     1923
population          3879
households          1809
median_income       12825
median_house_value  3833
ocean_proximity      5
dtype: int64
```

since there is no one feild that has unique values for the whole dataframe we will create a primary key to enabling merging of datasets.

In [5]:

```
ey = 1
for o, row in housing_raw.iterrows():
    housing_raw.loc[o, 'id'] = ey
    ey += 1
```

C:\Users\benjo\AppData\Local\Temp\ipykernel_4548\1766678878.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
housing_raw.loc[o, 'id'] = ey
```

The cell below creates 3 new columns to help us search for outliers

In [6]:

```
housing_raw['population_density'] = housing_raw['population']/housing_raw['households']
housing_raw['bedless'] = housing_raw['population']/housing_raw['total_bedrooms']
housing_raw.describe()
```

C:\Users\benjo\AppData\Local\Temp\ipykernel_4548\4209021603.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
housing_raw['population_density'] = housing_raw['population']/housing_raw['households']
```

C:\Users\benjo\AppData\Local\Temp\ipykernel_4548\4209021603.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
housing_raw['bedless'] = housing_raw['population']/housing_raw['total_bedrooms']
```

Out[6]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	p
count	20433.000000	20433.000000	20433.000000	20433.000000	20433.000000	20433.000000
mean	-119.570689	35.633221	28.633094	2636.504233	537.870553	142.000000
std	2.003578	2.136348	12.591805	2185.269567	421.385070	113.000000
min	-124.350000	32.540000	1.000000	2.000000	1.000000	0.000000
25%	-121.800000	33.930000	18.000000	1450.000000	296.000000	78.000000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	116.000000
75%	-118.010000	37.720000	37.000000	3143.000000	647.000000	172.000000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	3568.000000

keeping valid data


In []:

In []:

Plotting box plots to show outliers

In [215]:

```
sns.boxplot(x=housing_raw['housing_median_age'])
plt.show()
sns.boxplot(x=housing_raw['total_rooms'])
plt.show()
sns.boxplot(x=housing_raw['total_bedrooms'])
plt.show()
sns.boxplot(x=housing_raw['population'])
plt.show()
sns.boxplot(x=housing_raw['households'])
plt.show()
sns.boxplot(x=housing_raw['median_income'])
plt.show()
sns.boxplot(x=housing_raw['median_house_value'])
plt.show()
sns.boxplot(x=housing_raw['population_density'])
plt.show()
sns.boxplot(x=housing_raw['bedless'])
plt.show()
```



In [7]:

```
#calculate upper fences
```

```
wisker_fence_population = np.percentile(housing_raw['population'], 75) + 1.5 * (np.percentile(housing_raw['population'], 75) - np.percentile(housing_raw['population'], 25))  
upper_fence_population = wisker_fence_population + 1.5 * wisker_fence_population
```

```
wisker_fence_total_bedrooms = np.percentile(housing_raw['total_bedrooms'], 75) + 1.5 * (np.percentile(housing_raw['total_bedrooms'], 75) - np.percentile(housing_raw['total_bedrooms'], 25))  
upper_fence_total_bedrooms = wisker_fence_total_bedrooms + 1.5 * wisker_fence_total_bedrooms
```

```
wisker_fence_households = np.percentile(housing_raw['households'], 75) + 1.5 * (np.percentile(housing_raw['households'], 75) - np.percentile(housing_raw['households'], 25))  
upper_fence_households = wisker_fence_households + 1.5 * wisker_fence_households
```

```
wisker_fence_median_income = np.percentile(housing_raw['median_income'], 75) + 1.5 * (np.percentile(housing_raw['median_income'], 75) - np.percentile(housing_raw['median_income'], 25))  
upper_fence_median_income = wisker_fence_median_income + 1.5 * wisker_fence_median_income
```

```
wisker_fence_median_income = np.percentile(housing_raw['bedless'], 75) + 1.5 * (np.percentile(housing_raw['bedless'], 75) - np.percentile(housing_raw['bedless'], 25))  
upper_fence_median_income = wisker_fence_median_income + 1.5 * wisker_fence_median_income
```


```
wisker_fence_median_income = np.percentile(housing_raw['population_density'], 75) + 1.5 * (np.percentile(housing_raw['population_density'], 75) - np.percentile(housing_raw['population_density'], 25))  
upper_fence_median_income = wisker_fence_median_income + 1.5 * wisker_fence_median_income
```

```
# Filter outliers
```

```
housing = housing_raw[(housing_raw['population'] <= upper_fence_population)]  
housing = housing[(housing['total_bedrooms'] <= upper_fence_total_bedrooms)]  
housing = housing[(housing['households'] <= upper_fence_households)]  
housing = housing[(housing['median_income'] <= upper_fence_median_income)]  
housing = housing[(housing['bedless'] <= upper_fence_median_income)]  
housing = housing[(housing['population_density'] <= upper_fence_median_income)]
```



```
#plot boxes again
```

```
sns.boxplot(x=housing['total_rooms'])  
plt.show()  
sns.boxplot(x=housing['total_bedrooms'])  
plt.show()  
sns.boxplot(x=housing['population'])  
plt.show()  
sns.boxplot(x=housing['households'])  
plt.show()  
sns.boxplot(x=housing['median_income'])  
plt.show()  
sns.boxplot(x=housing['population_density'])  
plt.show()  
sns.boxplot(x=housing['bedless'])  
plt.show()
```



In [188]:

```
sns.boxplot(x=housing['population_density'])  
plt.show()  
sns.boxplot(x=housing['population'])  
plt.show()
```



In [10]:

```
housing.describe()
```

Out[10]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	p
count	20167.000000	20167.000000	20167.000000	20167.000000	20167.000000	20167.000000
mean	-119.572935	35.637605	28.694055	2563.436406	525.327862	1381.214763
std	2.003969	2.138954	12.527329	1846.640076	362.369160	98.461560
min	-124.350000	32.540000	1.000000	2.000000	2.000000	1.000000
25%	-121.800000	33.930000	18.000000	1449.500000	297.000000	78.000000
50%	-118.500000	34.260000	29.000000	2120.000000	435.000000	116.000000
75%	-118.010000	37.720000	37.000000	3120.000000	644.000000	176.000000
max	-114.310000	41.950000	52.000000	18634.000000	2885.000000	771.000000

As we can see, significant outliers have been taken out, yet we still remain below our 5% removal limit so we will now focus on cleaning the capped 500000

we found that houses that are on islands are always very expensive, therefore any rows that have 500000 as average house value will be kept we found that houses close to the city are also tend to be very expensive so any rows with houses close to the city will be kept

first we need to know the distance from the city

In [8]:

```

import math

housing.loc[:, 'distance_to_closest_city'] = 0

for r, block_row in housing.iterrows(): #for each row
    lati = block_row['latitude'] #create variable to hold latitude
    longi = block_row['longitude'] #create variable to hold longitude
    latcitydif = [] #empty a list hold the differences between the block & citys
    longcitydif = []
    distance = [] #empty a list to distances between block and city's

    for p, city_row in california_cities.iterrows(): #for each row in cities df
        # Subtract the latitude value from each row in the second dataframe
        latdiff = lati - city_row['latitude']
        longdiff = longi - city_row['longitude']
        # Append the resulting series to the list
        latcitydif.append(latdiff)
        longcitydif.append(longdiff)

    for d in range(len(latcitydif)): #for every value in list
        #find the distance using the euclidean distance
        sq = (latcitydif[d])**2 + (longcitydif[d])**2
        dist = math.sqrt(sq)
        #append the distance to distance list
        distance.append(dist)
    #identify the minimum distance in the list
    min_distance = min(distance)
    #append minimum in list to column
    housing.loc[r, 'distance_to_closest_city'] = min_distance

housing.describe()

```

Out[8]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	p
count	20167.000000	20167.000000	20167.000000	20167.000000	20167.000000	20167.000000
mean	-119.572935	35.637605	28.694055	2563.436406	525.327862	138.015500
std	2.003969	2.138954	12.527329	1846.640076	362.369160	98.989100
min	-124.350000	32.540000	1.000000	2.000000	2.000000	1.000000
25%	-121.800000	33.930000	18.000000	1449.500000	297.000000	78.000000
50%	-118.500000	34.260000	29.000000	2120.000000	435.000000	118.000000
75%	-118.010000	37.720000	37.000000	3120.000000	644.000000	177.000000
max	-114.310000	41.950000	52.000000	18634.000000	2885.000000	777.000000


In [11]:

```
#create dataframes consisting of island houses and
df_city_houses = housing[(housing['distance_to_closest_city'] <= 0.15)]
df_island_houses = housing[(housing['ocean_proximity'] == 'island')]

safe_rows = pd.concat([df_island_houses, df_city_houses], ignore_index=True, sort=False)
df_island_city = safe_rows.describe()

housing1 = housing[(housing['median_house_value'] <= 500000)]
housing2 = pd.concat([housing1, safe_rows], ignore_index=True, sort=False)
housing = housing2.drop_duplicates()
housing2.describe()

sns.boxplot(x=df_city_houses['total_rooms'])
plt.show()
sns.boxplot(x=df_city_houses['total_bedrooms'])
plt.show()
sns.boxplot(x=df_city_houses['population'])
plt.show()
sns.boxplot(x=df_city_houses['households'])
plt.show()
sns.boxplot(x=df_city_houses['median_income'])
plt.show()
sns.boxplot(x=df_city_houses['population_density'])
plt.show()
sns.boxplot(x=df_city_houses['bedless'])
plt.show()
```



In [12]:

```
housing.describe()  
#note, original count was 20433, now it's 19506, 972 rows taken out, approximately 4.5%
```

Out[12]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	p
count	19543.000000	19543.000000	19543.000000	19543.000000	19543.000000	19543.000000
mean	-119.575704	35.657503	28.623343	2548.926828	526.160927	136.135612
std	2.008108	2.148433	12.544141	1829.788661	361.904789	96.030914
min	-124.350000	32.540000	1.000000	2.000000	2.000000	1.000000
25%	-121.780000	33.930000	18.000000	1444.000000	298.000000	78.000000
50%	-118.510000	34.270000	29.000000	2111.000000	436.000000	117.000000
75%	-118.000000	37.730000	37.000000	3109.000000	645.000000	172.000000
max	-114.310000	41.950000	52.000000	18634.000000	2885.000000	771.000000

The code below creates a new column with the distance a block is from a major city

In [241]:

Out[241]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	p
count	20258.000000	20258.000000	20258.000000	20258.000000	20258.000000	20258.000000
mean	-119.575046	35.636977	28.754418	2537.309409	520.139599	137.135612
std	2.003614	2.137204	12.523427	1750.309365	351.709278	96.030914
min	-124.350000	32.540000	1.000000	2.000000	2.000000	1.000000
25%	-121.800000	33.930000	18.000000	1448.000000	295.000000	78.000000
50%	-118.500000	34.260000	29.000000	2119.000000	433.000000	117.000000
75%	-118.010000	37.720000	37.000000	3114.000000	641.000000	172.000000
max	-114.310000	41.950000	52.000000	14125.000000	2823.000000	761.000000

The cell below creates some new dataframes containing the data in the area of los angeles and san francisco

In [28]:

```

longlostangel = housing[ (housing['longitude'] <= -117) & (housing['longitude'] >= -119)
lostangel = longlostangel[ (longlostangel['latitude'] >= 33) & (longlostangel['latitude']
santa = housing[ (housing['longitude'] <= -121) & (housing['longitude'] >= -123)]
san = santa[ (santa['latitude'] >= 37) & (santa['latitude'] <= 39)]
san.head(10)

```

Out[28]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	househ
0	-122.23	37.88	41	880	129.0	322	
1	-122.22	37.86	21	7099	1106.0	2401	
2	-122.24	37.85	52	1467	190.0	496	
3	-122.25	37.85	52	1274	235.0	558	
4	-122.25	37.85	52	1627	280.0	565	
5	-122.25	37.85	52	919	213.0	413	
6	-122.25	37.84	52	2535	489.0	1094	
7	-122.25	37.84	52	3104	687.0	1157	
8	-122.26	37.84	42	2555	665.0	1206	
9	-122.25	37.84	52	3549	707.0	1551	

Bedlessness

Import shapefile

In [13]:

```

from shapely.geometry import Point, Polygon
import os
import seaborn as sns
import shapely.ops
import netCDF4
from pyproj import Transformer
import matplotlib.patches as mpatches
import matplotlib.lines as mlines
import geopandas as gpd

```

In [20]:

```

os.environ['SHAPE_RESTORE_SHX'] = 'YES' #for some reason I got an error that told me to
#importing our stuff
county_shapefile = gpd.read_file('CA_Counties_TIGER2016.shp') #shape file for Cali map
county_shapefile['geometry'] = county_shapefile['geometry'].apply(lambda geom: shapely.o


```

Regression Plot of population to bedrooms

In [144]:

```
sns.regplot(data=housing, x=housing['population'], y=housing['total_bedrooms'],line_kws=
plt.title('not enough rooms for everyone')
plt.show()
p1 = np.poly1d(np.polyfit(housing['population'], housing['total_bedrooms'], 1))
pv = p1(housing['population'])
r2 = r2_score(housing['total_bedrooms'],pv)
print(p1)
print(r2)
```

#finding the % of people that have to share a room



0.3288 x + 65.78
0.7481271644738751

In [74]:

```
housing.loc[:, 'alone'] = housing['bedless'] < 2
sharing = housing['alone'].value_counts()[False]
not_sharing = housing['alone'].value_counts()[True]
sharing / (not_sharing + sharing)
```

C:\Users\benjo\AppData\Local\Temp\ipykernel_23440\1235303491.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
housing.loc[:, 'alone'] = housing['bedless'] < 2
```

Out[74]:

0.8569749015999591

On average 70% of Californians either share a bedroom or don't have a bedroom to sleep in. A significant amount of these bedless people reside in the Los Angeles area and the San Francisco area as seen in the map below showing the locations where the population and total number of bedrooms exceeds 4:1. We can also see from the regression plot showing the relationship between distance from major city and homelessness, that bedlessness is more frequent the closer you get to a city

In [220]:

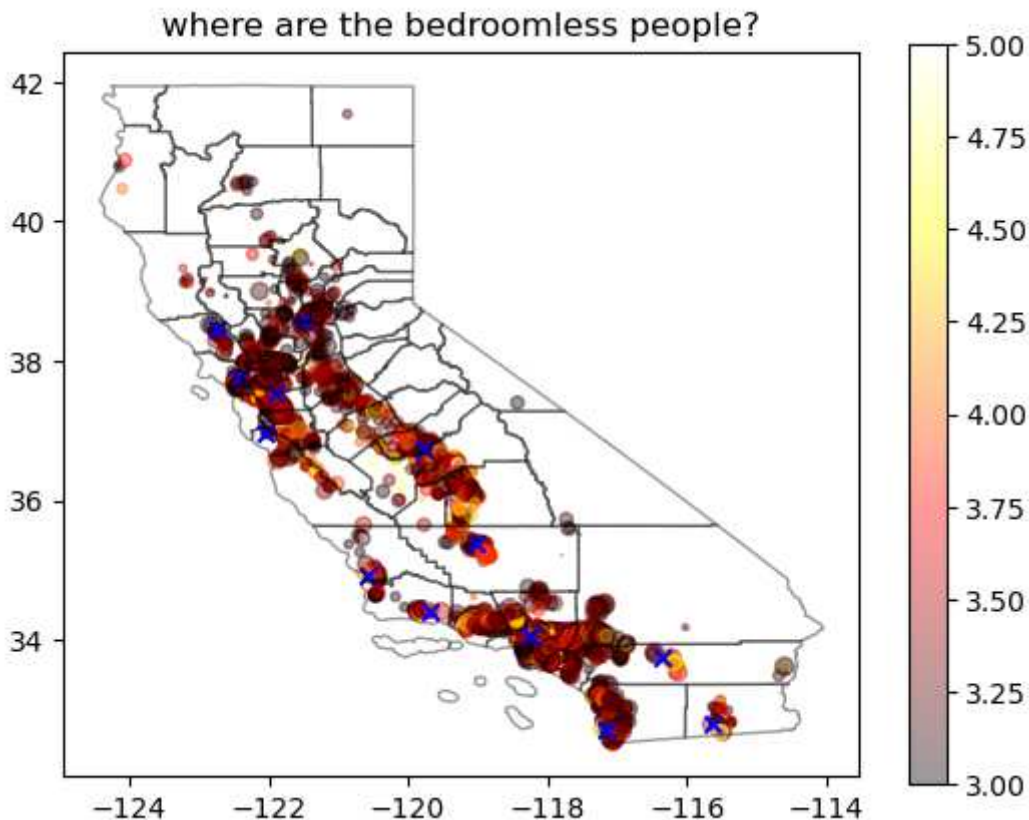
```

calirooms = housing[(housing['population'] >= 3*housing['total_bedrooms'])]
#make sure there are no outliers so that we can visualise better:
upper_fence_rooms_total_bedrooms = np.percentile(calirooms['total_bedrooms'], 75) + 1.5
upper_fence_rooms_population = np.percentile(calirooms['population'], 75) + 1.5 * (np.pe
lower_fence_rooms_total_bedrooms = np.percentile(calirooms['total_bedrooms'], 25) - 1.5
lower_fence_rooms_population = np.percentile(calirooms['population'], 25) - 1.5 * (np.pe
calirooms = calirooms[(calirooms['population'] >= lower_fence_rooms_population) &(calirooms

#lets set the coloring
calirooms['coloring'] = 0
for m, row in calirooms.iterrows():
    if row['population'] / row['total_bedrooms'] <= 5:
        calirooms.at[m, 'coloring'] = row['population'] / row['total_bedrooms']
    else:
        calirooms.at[m, 'coloring'] = 5

#now lets plot it with respect to a map
fig, ax = plt.subplots()
plt.scatter(x=calirooms['longitude'],y=calirooms['latitude'],alpha=0.4,s=50*calirooms['h
plt.colorbar()
plt.scatter(x=california_cities['longitude'],y=california_cities['latitude'],alpha = 1,m
#for i in range(california_cities.shape[0]):
#    plt.text(x=california_cities.longitude[i]+0.1,y=california_cities.latitude[i]+0.1,s
county_shapefile.plot(ax=ax, color='None', edgecolor='black',alpha=0.4)
plt.title('where are the bedroomless people?')
plt.show();

```



In [140]:

```

calirooms = san[(san['population'] >= 3*san['total_bedrooms'])]

#make sure there are no outliers so that we can visualise better:
#upper_fence_rooms_total_bedrooms = np.percentile(calirooms['total_bedrooms'], 75) + 1.5
#upper_fence_rooms_population = np.percentile(calirooms['population'], 75) + 1.5 * (np.p
#lower_fence_rooms_total_bedrooms = np.percentile(calirooms['total_bedrooms'], 25) - 1.5
#lower_fence_rooms_population = np.percentile(calirooms['population'], 25) - 1.5 * (np.p
#calirooms = calirooms[(calirooms['population'] >= lower_fence_rooms_population) &(calir

#lets set color
calirooms['coloring'] = 0
for m, row in calirooms.iterrows():
    if row['population'] / row['total_bedrooms'] <= 5:
        calirooms.at[m, 'coloring'] = row['population'] / row['total_bedrooms']
    else:
        calirooms.at[m, 'coloring'] = 5

#now lets plot it with respect to a map
fig, ax = plt.subplots()

plt.scatter(x=calirooms['longitude'],y=calirooms['latitude'],alpha=0.25,c=calirooms['col
#plt.scatter(x=calirooms_large['longitude'],y=calirooms_large['latitude'],alpha=0.5,c=ca
plt.colorbar()
plt.title('San Francisco & Santa Cruz area')
#plt.plot(-122.030797,37.090017, marker="o", markersize=25, markeredgecolor="blue", mark
#plt.text(-122.230797,37.154117, 'Santa Cruz',c='blue')
plt.plot(-121.889096, 37.354960, marker="o", markersize=50, markeredgecolor="blue", mark
plt.text(-121.759096, 37.554960, 'San Jose',c='blue')
plt.plot(-122.469417,37.704931, marker="o", markersize=40, markeredgecolor="blue", marke
plt.text(-122.856034,37.883974, 'San Francisco',c='blue')
plt.plot(-121.3944,38.581572, marker="o", markersize=50, markeredgecolor="blue", markerf
plt.text(-121.3244,38.801572, 'Sacramento',c='blue')

plt.show();

```

```
C:\Users\benjo\AppData\Local\Temp\ipykernel_23440\273978217.py:11: Setting
WithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
calirooms['coloring'] = 0
```

```
C:\Users\benjo\AppData\Local\Temp\ipykernel_23440\273978217.py:14: Setting
WithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
calirooms.at[m, 'coloring'] = row['population'] / row['total_bedrooms']
```



In [222]:


```

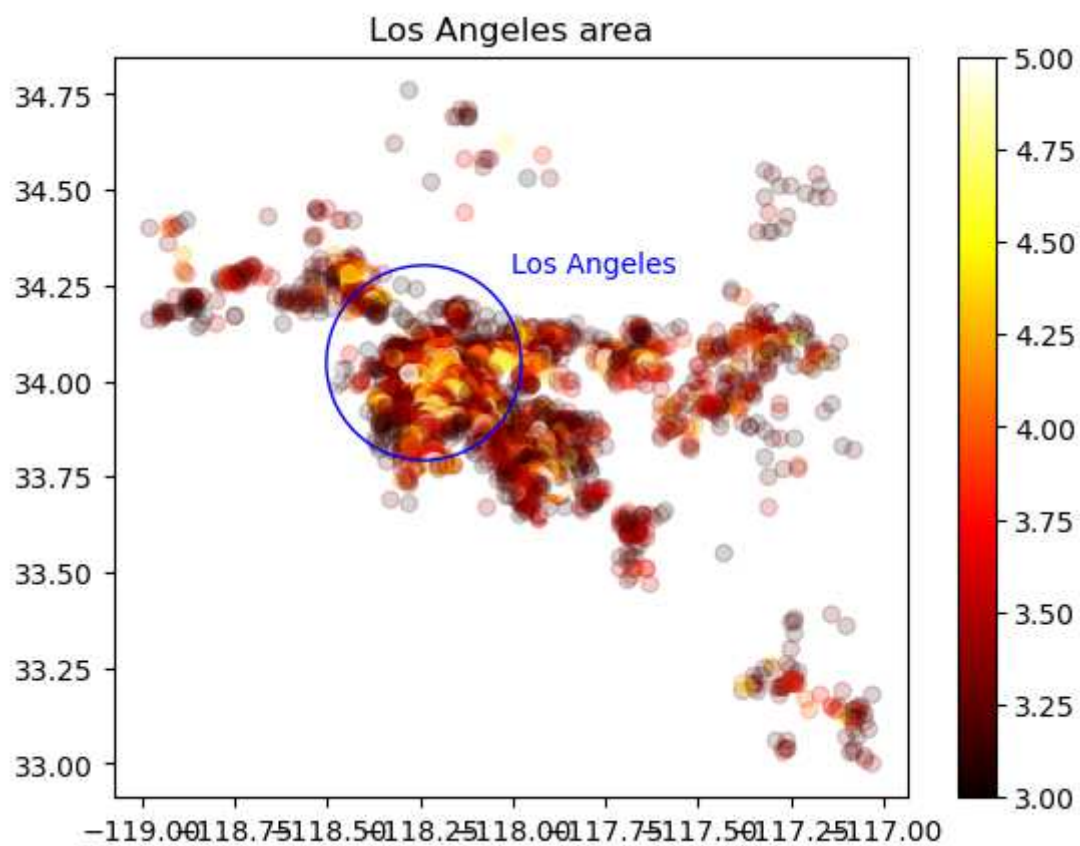
calirooms = lostangel[(lostangel['population'] >= 3*lostangel['total_bedrooms'])]
#make sure there are no outliers so that we can visualise better:
upper_fence_rooms_total_bedrooms = np.percentile(calirooms['total_bedrooms'], 75) + 1.5
upper_fence_rooms_population = np.percentile(calirooms['population'], 75) + 1.5 * (np.pe
lower_fence_rooms_total_bedrooms = np.percentile(calirooms['total_bedrooms'], 25) - 1.5
lower_fence_rooms_population = np.percentile(calirooms['population'], 25) - 1.5 * (np.pe
calirooms = calirooms[(calirooms['population'] >= lower_fence_rooms_population) & (calirooms['total_bedrooms'] >= lower_fence_rooms_total_bedrooms)]
#lets set color
calirooms['coloring'] = 0
for m, row in calirooms.iterrows():
    if row['population'] / row['total_bedrooms'] <= 5:
        calirooms.at[m, 'coloring'] = row['population'] / row['total_bedrooms']
    else:
        calirooms.at[m, 'coloring'] = 5

#now lets plot it with respect to a map
fig, ax = plt.subplots()

#now lets plot it with respect to a map
fig, ax = plt.subplots()
plt.scatter(x=calirooms['longitude'],y=calirooms['latitude'],alpha=calirooms['bedless']/
#plt.scatter(x=calirooms['longitude'],y=calirooms['latitude'],alpha=0.5,c=calirooms['pop
plt.colorbar()
plt.legend()
plt.title('Los Angeles area')
plt.plot(-118.243686,34.052233, marker="o", markersize=70, markeredgecolor="blue", marke
plt.text(-118.005,34.282233, 'Los Angeles',c='blue')
plt.show();
lostcalirooms = calirooms

```






In [202]:

```
plt.title('bedlessness in relation to proximity to city')
p2 = np.poly1d(np.polyfit(housing['distance_to_closest_city'], housing['bedless'], 2))
housefit = np.linspace(housing['distance_to_closest_city'].min(), housing['distance_to_c
plt.scatter(x=housing['distance_to_closest_city'], y=housing['bedless'],alpha=0.5,c=hous
plt.plot(housefit,p2(housefit),label="Average",c='green')
pv = p2(housing['distance_to_closest_city'])
r2 = r2_score(housing['bedless'],pv)
print(r2)
plt.show()


plt.title('bedlessness in los angeles in relation to proximity to los angeles')
p2 = np.poly1d(np.polyfit(lostangel['distance_to_closest_city'], lostangel['bedless'], 2))
housefit = np.linspace(lostangel['distance_to_closest_city'].min(), lostangel['distance_
plt.scatter(x=lostangel['distance_to_closest_city'], y=lostangel['bedless'],alpha=0.5,c=
plt.plot(housefit,p2(housefit),label="Average",c='red')
pv = p2(lostangel['distance_to_closest_city'])
r2 = r2_score(lostangel['bedless'],pv)
print(r2)
plt.show()

plt.title('bedlessness in San Fransisco in relation to proximity to San Francisco, San J
p2 = np.poly1d(np.polyfit(san['distance_to_closest_city'], san['bedless'], 2))
housefit = np.linspace(san['distance_to_closest_city'].min(), san['distance_to_closest_c
plt.scatter(x=san['distance_to_closest_city'], y=san['bedless'],alpha=0.5,c=san['bedless
plt.plot(housefit,p2(housefit),label="Average",c='red')
pv = p2(san['distance_to_closest_city'])
r2 = r2_score(san['bedless'],pv)
print(r2)
plt.show()
```

0.016053007143745024



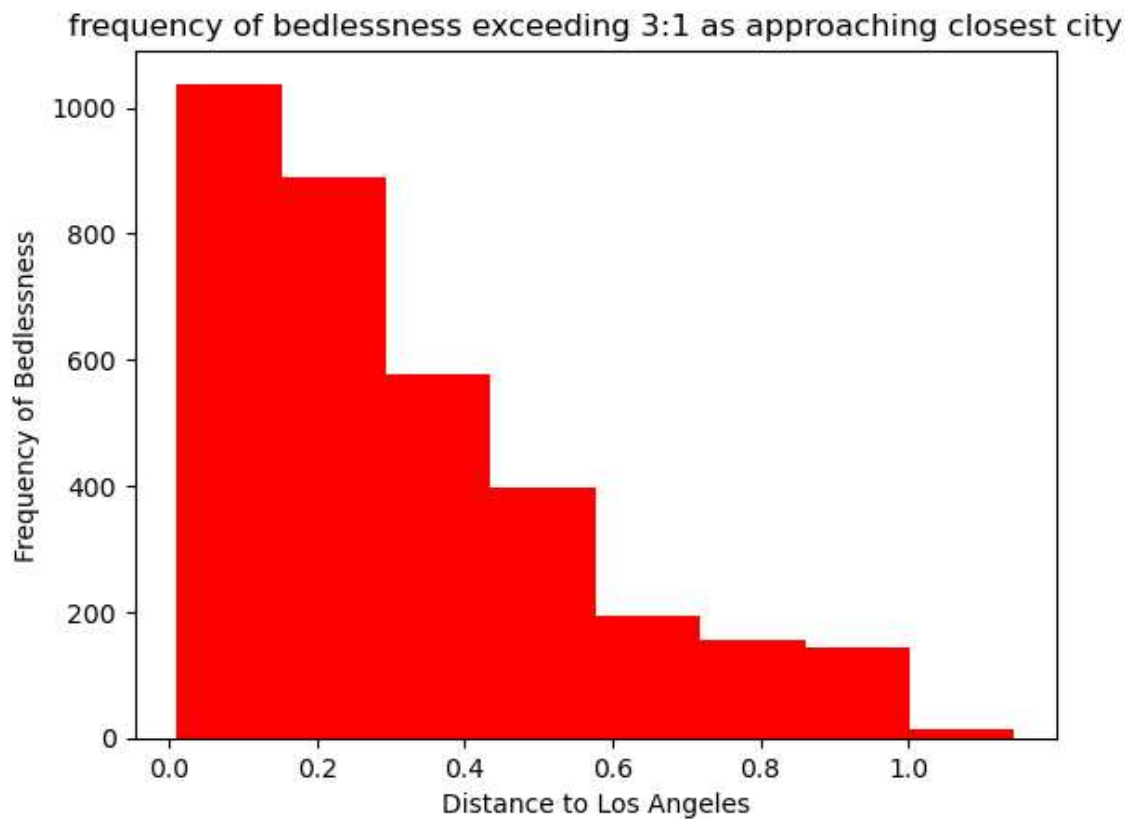
0.011681415152787755



0.012205883191009459

bedlessness in San Fransisco in relation to proximity to San Francisco, San Jose and Santa Cruz


```
lostcalrooms.hist('distance_to_closest_city',bins=8,grid=False,color='red')  
plt.title('frequency of bedlessness exceeding 3:1 as approaching closest city')  
plt.xlabel('Distance to Los Angeles')  
plt.ylabel('Frequency of Bedlessness')  
plt.show()
```



In [206]:

```
#plotting frequency
```

```
calirooms.hist('distance_to_closest_city',bins=8,grid=False,color='red')  
plt.title('frequency of bedlessness exceeding 3:1 as approaching closest city')  
plt.xlabel('Distance to Los Angeles')  
plt.ylabel('Frequency of Bedlessness')  
plt.show()
```



Median income and house value as approaching city

Median Income

In [21]:

```
caliincomebin = pd.cut(housing['median_income'], 4, precision=2)
calivalubin = pd.cut(housing['median_house_value'], 4, precision=2)

conditions = [
    (housing['median_income'] <= 2.24),
    (housing['median_income'] > 2.24) & (housing['median_income'] <= 4.48),
    (housing['median_income'] > 4.48) & (housing['median_income'] <= 8),
    (housing['median_income'] > 8)
]
fence_house_value = np.percentile(housing['median_house_value'], 75) + 1.5 * (np.percent
Vconditions = [
    (housing['median_house_value'] <= np.percentile(housing['median_house_value'], 25)),
    (housing['median_house_value'] > np.percentile(housing['median_house_value'], 25)) &
    (housing['median_house_value'] > np.percentile(housing['median_house_value'], 75)) &
    (housing['median_house_value'] > fence_house_value)
]
conditions_age = [
    (housing['housing_median_age'] <= 12),
    (housing['housing_median_age'] > 10) & (housing['housing_median_age'] <= 30),
    (housing['housing_median_age'] > 30) & (housing['housing_median_age'] <= 50),
    (housing['housing_median_age'] > 50)
]

status = ['poor', 'middle_class', 'wealthy', 'uber_rich']
housing['status'] = np.select(conditions, status)
house_bins = ['Low_Value', 'Medium_Value', 'High_Value', 'Highest_Value']
housing['house_bin'] = np.select(Vconditions, house_bins)
housing.head()
age_bins = ['<10', '10<30', '30<50', '50<']
housing['age_bin'] = np.select(conditions_age, age_bins)
```

```
C:\Users\benjo\AppData\Local\Temp\ipykernel_23440\450019214.py:25: Setting
WithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
housing['status'] = np.select(conditions, status)
```

```
C:\Users\benjo\AppData\Local\Temp\ipykernel_23440\450019214.py:27: Setting
WithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
housing['house_bin'] = np.select(Vconditions, house_bins)
```

```
C:\Users\benjo\AppData\Local\Temp\ipykernel_23440\450019214.py:30: Setting
WithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
housing['age_bin'] = np.select(conditions_age, age_bins)
```

In [173]:

```
#california
fig, ax = plt.subplots()

status_colors = {'poor': 'black', 'middle_class': 'red', 'wealthy': 'orange', 'uber_rich': 'blue'}

# Map the categories to colors for each data point
statuscolors = housing['status'].map(status_colors)

#county_shapefile.plot(ax=ax, color='white', edgecolor='black')

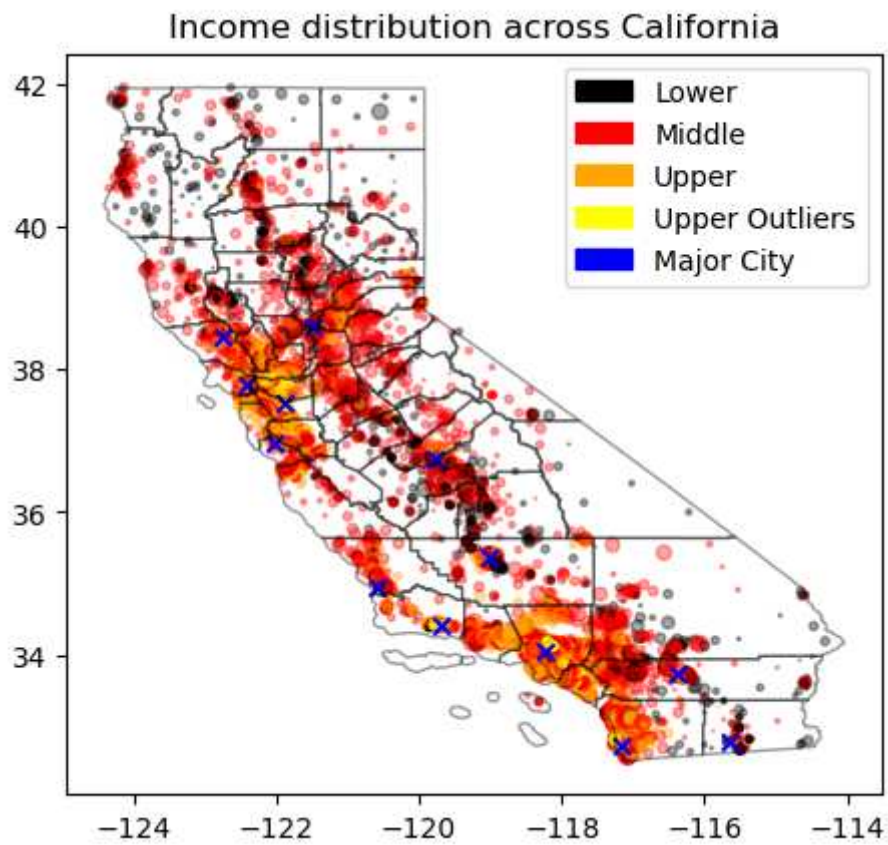
plt.scatter(x=housing['longitude'],y=housing['latitude'],alpha = (35+((1.066**((100*(housing['income']-10000)/(100000-10000))))))

plt.scatter(x=california_cities['longitude'],y=california_cities['latitude'],alpha = 1,marker='o')
#for i in range(california_cities.shape[0]):
#    plt.text(x=california_cities.longitude[i]+0.1,y=california_cities.latitude[i]+0.1,s=f'{california_cities.city[i]}',color='black',fontSize=8)
county_shapefile.plot(ax=ax, color='None', edgecolor='black',alpha=0.4)

plt.title("Income distribution across California")

outliers = mpatches.Patch(color='yellow', label='Upper Outliers')
poor = mpatches.Patch(color='black', label='Lower')
mid = mpatches.Patch(color='red', label='Middle')
upper = mpatches.Patch(color='orange', label='Upper')
city = mpatches.Patch(color='blue', label='Upper')
plt.legend(handles=[poor,mid,upper,outliers,city])

plt.show();
```

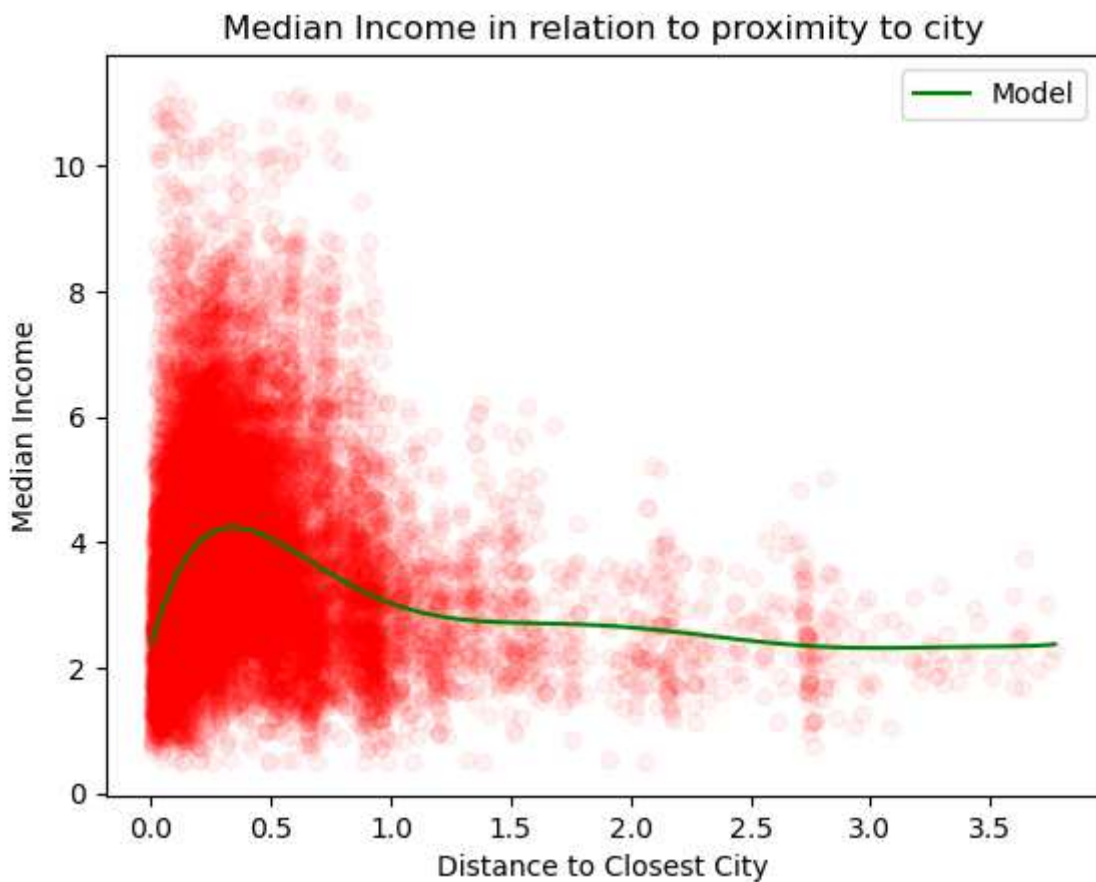


In []:

In [195]:

```
plt.title('Median Income in relation to proximity to city')
p2 = np.poly1d(np.polyfit(housing['distance_to_closest_city'], housing['median_income'],
housefit = np.linspace(housing['distance_to_closest_city'].min(), housing['distance_to_c
plt.scatter(x=housing['distance_to_closest_city'], y=housing['median_income'], c='red',
plt.plot(housefit, p2(housefit), label="Model", c='green')
print(r2)
plt.legend()
plt.xlabel('Distance to Closest City')
plt.ylabel('Median Income')
plt.show()
```

0.012205883191009459



In []:

Median House price

In [176]:

```

fig, ax = plt.subplots()

house_bin_colors = {'Low_Value': 'orange', 'Medium_Value': 'green', 'High_Value': 'blue'}

# Map the categories to colors for each data point
house_bin_colors = housing['house_bin'].map(house_bin_colors)

plt.scatter(x=housing['longitude'],y=housing['latitude'],alpha = (35+((1.066**((100*(housing['house_bin'] == 'Low_Value') - 1)))))

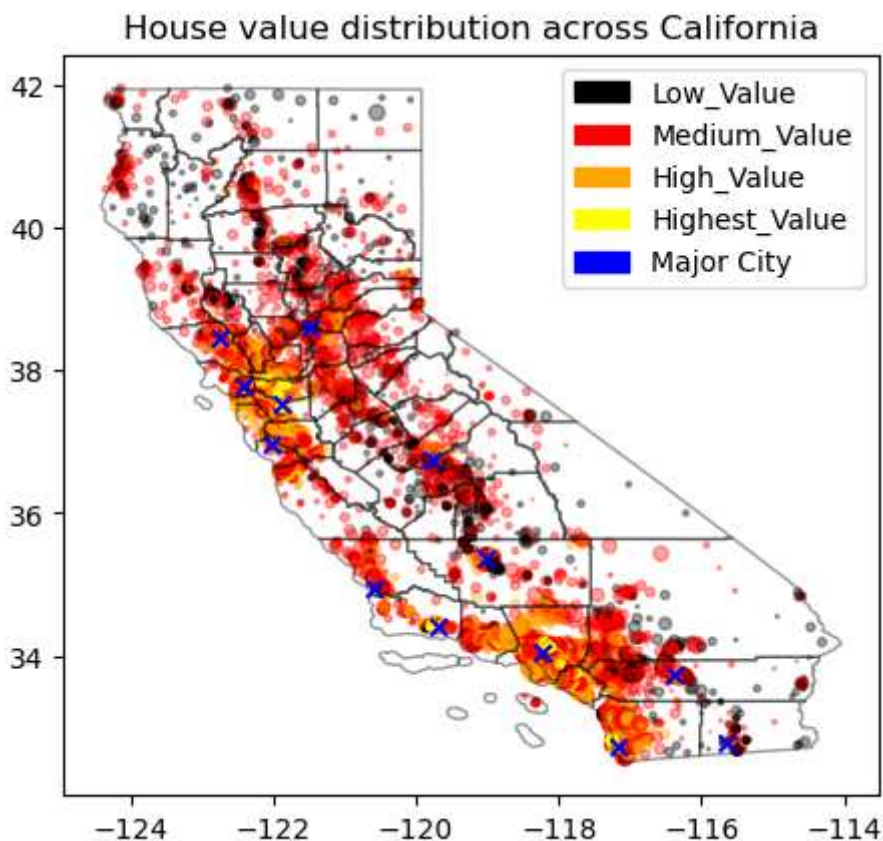
plt.title("House value distribution across California")

plt.scatter(x=california_cities['longitude'],y=california_cities['latitude'],alpha = 1,marker='x')

Highest_Value = mpatches.Patch(color='yellow', label='Highest_Value')
Low_Value = mpatches.Patch(color='black', label='Low_Value')
Medium_Value = mpatches.Patch(color='red', label='Medium_Value')
High_Value = mpatches.Patch(color='orange', label='High_Value')
city = mpatches.Patch(color='blue', label='Major City')
plt.legend(handles=[Low_Value,Medium_Value,High_Value,Highest_Value,city])

plt.show();

```

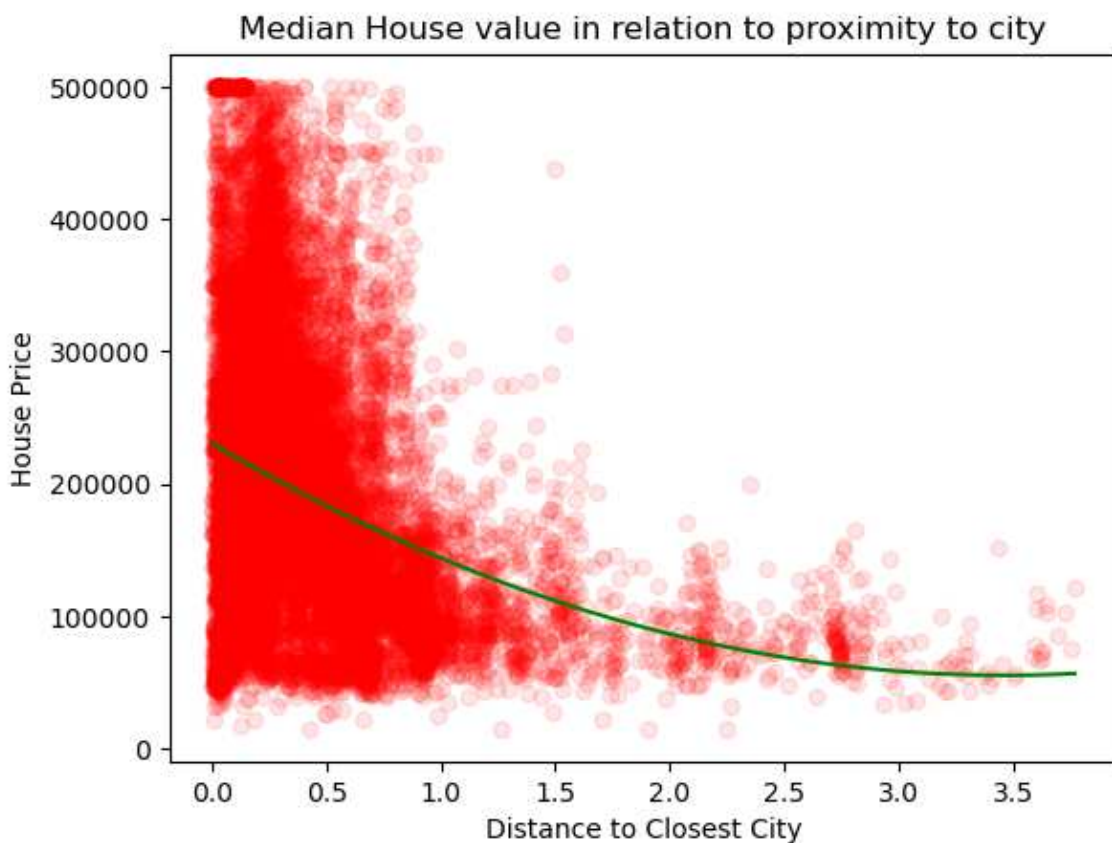


In []:

In [196]:

```
plt.title('Median House value in relation to proximity to city')
p2 = np.poly1d(np.polyfit(housing['distance_to_closest_city'], housing['median_house_val
housefit = np.linspace(housing['distance_to_closest_city'].min(), housing['distance_to_c
plt.scatter(x=housing['distance_to_closest_city'], y=housing['median_house_value'],c='re
plt.plot(housefit,p2(housefit),label="Average",c='green')
pv = p2(housing['distance_to_closest_city'])
r2 = r2_score(housing['median_house_value'],pv)
print(r2)
plt.xlabel('Distance to Closest City')
plt.ylabel('House Price')
plt.show()
```


0.09070381120347082



In [157]:

```
plt.title('Median House age in relation to proximity to city')
p2 = np.poly1d(np.polyfit(housing['distance_to_closest_city'], housing['housing_median_age'], 2))
housefit = np.linspace(housing['distance_to_closest_city'].min(), housing['distance_to_closest_city'].max(), 100)
plt.scatter(x=housing['distance_to_closest_city'], y=housing['housing_median_age'], c='red')
plt.plot(housefit, p2(housefit), label="Average", c='green')
pv = p2(housing['distance_to_closest_city'])
r2 = r2_score(housing['housing_median_age'], pv)
print(r2)
plt.show()
```

0.13657061055191722



In [152]:

```

fig, ax = plt.subplots()

age_bin_colors = {'<10': 'black', '10<30': 'red', '30<50': 'orange', '50<': 'yellow'}

# Map the categories to colors for each data point
age_bin_colors = housing['age_bin'].map(age_bin_colors)

plt.scatter(x=housing['longitude'],y=housing['latitude'],alpha = (35+((1.066**((100*(housing['age_bin'].map(age_bin_colors).get('age_bin')).get('age_bin')))))))

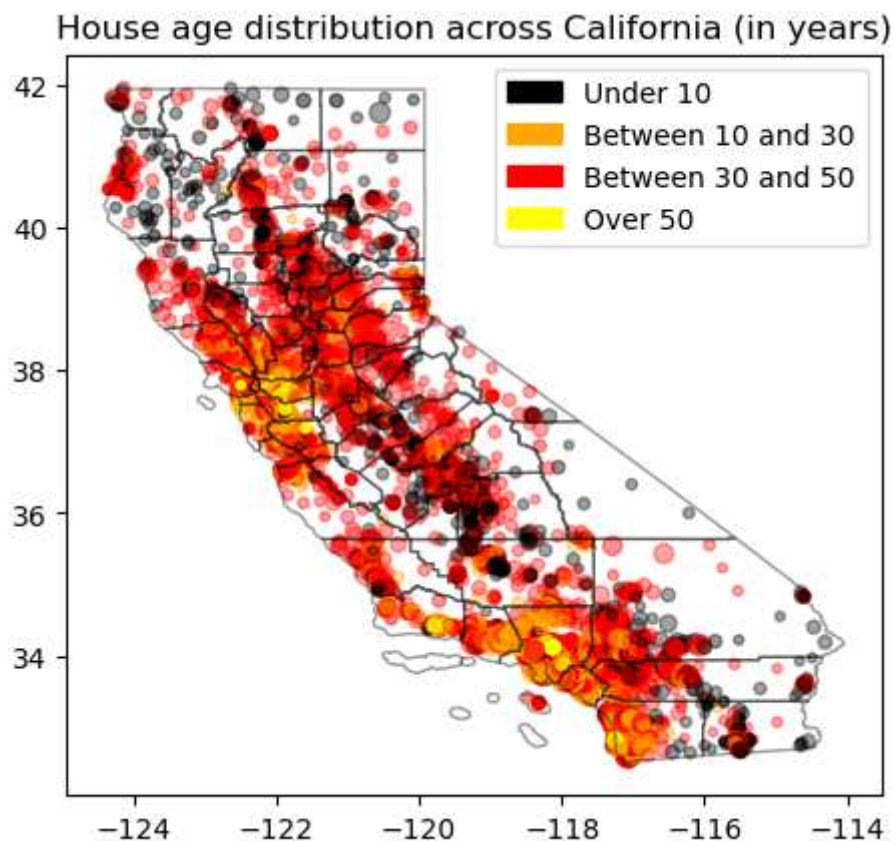
county_shapefile.plot(ax=ax, color='None', edgecolor='black',alpha=0.4)

plt.title("House age distribution across California (in years)")

outliers = mpatches.Patch(color='yellow', label='Over 50')
poor = mpatches.Patch(color='black', label='Under 10')
mid = mpatches.Patch(color='orange', label='Between 10 and 30')
upper = mpatches.Patch(color='red', label='Between 30 and 50')
plt.legend(handles=[poor,mid,upper,outliers])

plt.show();

```




In []:

In [15]:

```
sns.heatmap(housing.corr())
```

Out[15]:

<AxesSubplot:>




In [16]:

```
housing.hist(alpha=0.5, bins=50, edgecolor='black', figsize=(20,15))
```


Out[16]:

```
array([[<AxesSubplot:title={'center':'longitude'}>,
        <AxesSubplot:title={'center':'latitude'}>,
        <AxesSubplot:title={'center':'housing_median_age'}>],
       [<AxesSubplot:title={'center':'total_rooms'}>,
        <AxesSubplot:title={'center':'total_bedrooms'}>,
        <AxesSubplot:title={'center':'population'}>],
       [<AxesSubplot:title={'center':'households'}>,
        <AxesSubplot:title={'center':'median_income'}>,
        <AxesSubplot:title={'center':'median_house_value'}>],
       [<AxesSubplot:title={'center':'id'}>,
        <AxesSubplot:title={'center':'population_density'}>,
        <AxesSubplot:title={'center':'bedless'}>]], dtype=object)
```




In [19]:

```
plt.scatter(x=housing['longitude'],y=housing['latitude'])  
plt.show();  
#we can't see claerly
```



In [20]:

```
plt.scatter(x=housing['longitude'],y=housing['latitude'],alpha=0.1)  
plt.show();  
#we now can see the density of the blocks
```




In [21]:

```
plt.figure(figsize=(15,10))


plt.scatter(x=housing['longitude'],y=housing['latitude'],s=housing['population']/80,labels=housing['population'])
plt.colorbar()
plt.legend()
plt.show()
```

#the size of the circle represents the population and the colour of the circle represent



In [22]:


```
plt.figure(figsize=(15,10))  
  
plt.scatter(x=housing['longitude'],y=housing['latitude'],s=housing['population']/80,labels=housing['population'],c=housing['population']/80)  
plt.colorbar()  
plt.legend()  
plt.show()
```


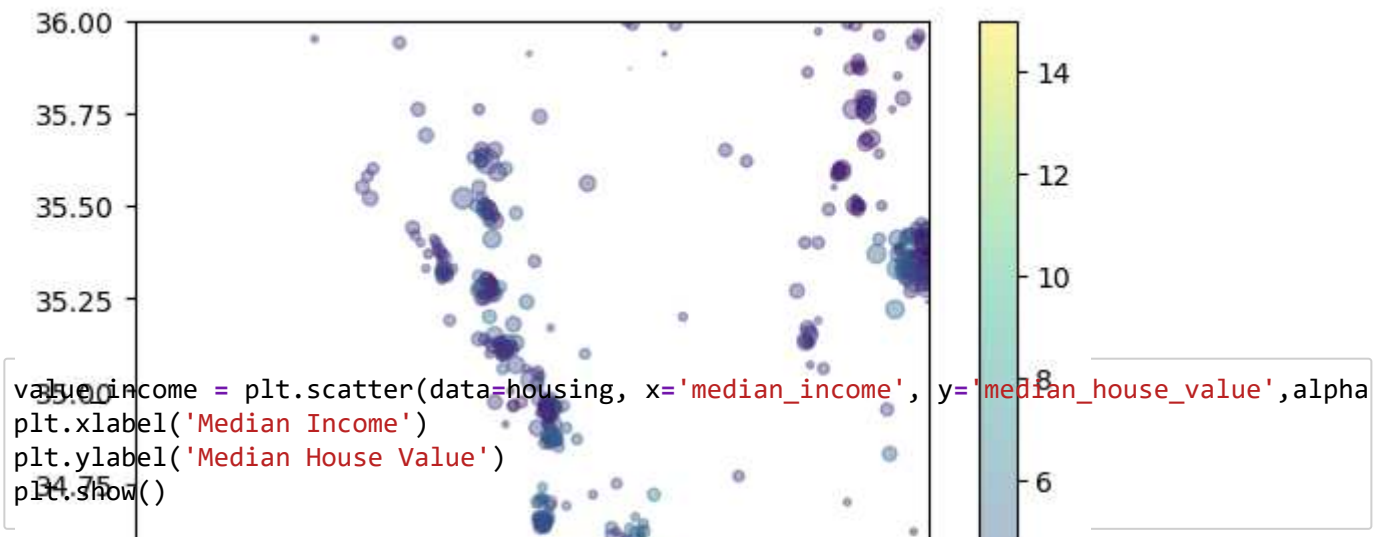


In [24]:

```
plt.scatter(x=housing['longitude'],y=housing['latitude'],s=housing['population']/80,labels=housing['population'])
plt.xlim([-122, -119])
plt.ylim([34, 36])
plt.colorbar()
plt.legend()
plt.show()

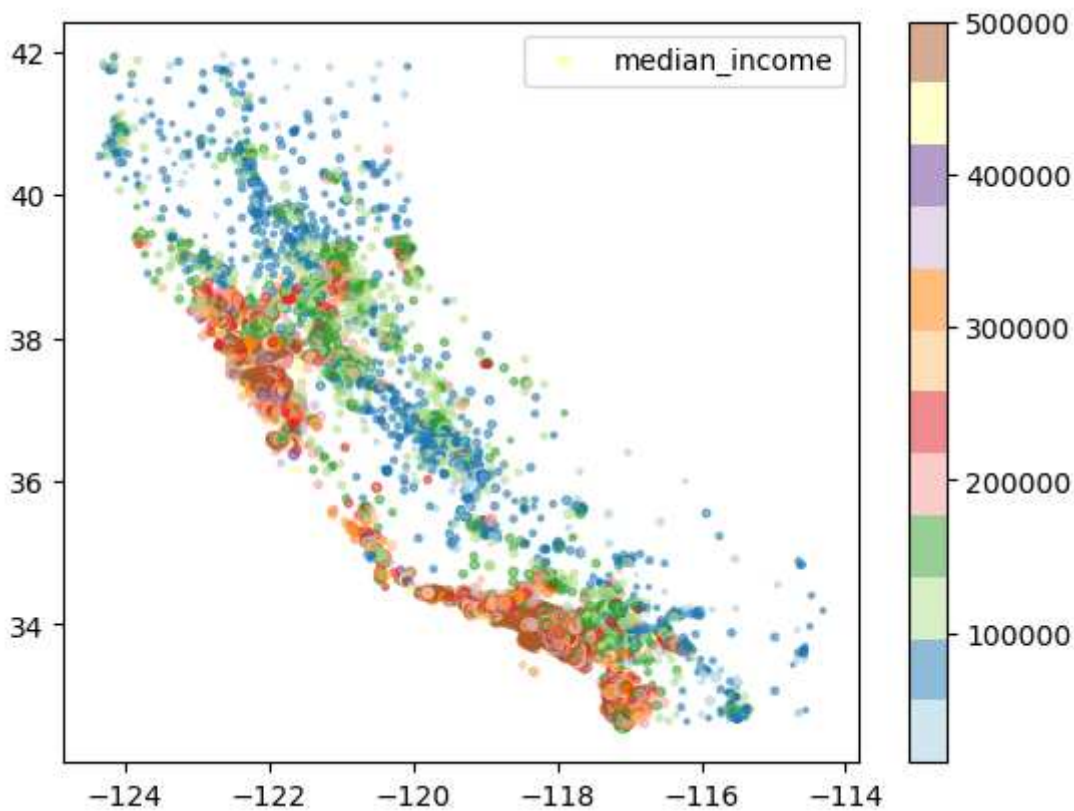
plt.scatter(x=housing['longitude'],y=housing['latitude'],s=housing['population']/80,labels=housing['population'])
plt.xlim([-122, -119])
plt.ylim([34, 36])
plt.colorbar()
plt.legend()
plt.show()
```





In [17]:

```
plt.scatter(x=housing['longitude'],y=housing['latitude'],s=housing['median_income']/0.5,
            label='median_income',alpha=0.5,c=housing['median_house_value'],cmap=plt.get_cmap('plasma'))
plt.colorbar()
plt.legend()
plt.show()
```



In [18]:

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score

X = housing[['median_income']]
Y = housing['median_house_value']

linear = LinearRegression(fit_intercept=True)
linear.fit(X, Y)
training_score = linear.score(X, Y)

print("Training score is",np.round(training_score, 3))
print("Correlation score is",np.round(np.sqrt(training_score), 3))
print("Coefficients are",np.round(linear.coef_, 3))
print("Intercept is",np.round(linear.intercept_,3))
```

Training score is 0.474
 Correlation score is 0.688
 Coefficients are [41928.972]
 Intercept is 44672.703

Population and Households Analysis

In [25]:

```
housing['population'].describe()
```

Out[25]:

```
count      20253.000000
mean       1374.551671
std        920.221690
min         3.000000
25%        785.000000
50%       1160.000000
75%       1703.000000
max       7694.000000
Name: population, dtype: float64
```

In [26]:

```
housing['households'].describe()
```

Out[26]:

```
count      20253.000000
mean        483.661137
std        319.975304
min         2.000000
25%        280.000000
50%        408.000000
75%        598.000000
max       2660.000000
Name: households, dtype: float64
```

In [28]:

```
X=pd.DataFrame(housing['households'])
Y=pd.DataFrame(housing['population'])
X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size=0.8,random_state=33)
reg = LinearRegression(fit_intercept=True).fit(X_train,Y_train)
training_score=reg.score(X_train,Y_train)
print('Regression coefficient estimate is',reg.coef_)
print('Regression y-int estimate is',reg.intercept_)
```

```
Regression coefficient estimate is [[2.67036084]]
Regression y-int estimate is [90.68257648]
```

In [30]:

```
training_score = reg.score(X_train,Y_train)
predictions = reg.predict(X_test)
testing_score = r2_score(Y_test,predictions)
print("Training score is",np.round(training_score, 3))
print("Testing score is",np.round(testing_score, 3))
```

```
Training score is 0.824
Testing score is 0.813
```

In [244]:

```

import math

housing.loc[:, 'distance_to_closest_city'] = 0

for r, block_row in housing.iterrows():
    lati = block_row['latitude']
    longi = block_row['longitude']
    latcitydif = []
    longcitydif = []
    distance = []

    for p, city_row in california_cities.iterrows():
        # Subtract the latitude value from each row in the second dataframe
        latdiff = lati - city_row['latitude']
        longdiff = longi - city_row['longitude']
        # Append the resulting series to the list
        latcitydif.append(latdiff)
        longcitydif.append(longdiff)

    for d in range(len(latcitydif)):
        sq = (latcitydif[d])**2 + (longcitydif[d])**2
        dist = math.sqrt(sq)
        distance.append(dist)

    min_distance = min(distance)
    housing.loc[r, 'distance_to_closest_city'] = min_distance

housing.describe()

```

Out[244]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	p
count	20253.000000	20253.000000	20253.000000	20253.000000	20253.000000	20253.000000
mean	-119.575193	35.636999	28.755098	2537.761072	520.231966	137.000000
std	2.003462	2.137220	12.524251	1750.246360	351.694421	92.000000
min	-124.350000	32.540000	1.000000	2.000000	2.000000	7.000000
25%	-121.800000	33.930000	18.000000	1448.000000	296.000000	78.000000
50%	-118.500000	34.260000	29.000000	2119.000000	433.000000	116.000000
75%	-118.010000	37.720000	37.000000	3114.000000	641.000000	170.000000
max	-114.310000	41.950000	52.000000	14125.000000	2823.000000	765.000000

In [229]:

```

#create dataframe containing highest income values
#create dataframe containing highest house value values
#inner merge both dataframes to a new dataframe, at end of outlier removal, merge dataframes

df_city_houses = housing[(housing['distance_to_closest_city'] <= 0.15)]
df_island_houses = housing[(housing['ocean_proximity'] == 'island')]

safe_rows = pd.concat([df_island_houses, df_city_houses], ignore_index=True, sort=False)
df_island_city = safe_rows.describe()

housing1 = housing[(housing['median_house_value'] <= 500000)]
housing2 = pd.concat([housing1, safe_rows], ignore_index=True, sort=False)
housing = housing2.drop_duplicates()
housing2.describe()

```

Out[229]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	pop
count	24244.000000	24244.000000	24244.000000	24244.000000	24244.000000	24244.000000
mean	-119.545886	35.589427	29.720714	2483.257672	523.736883	1385.199739
std	2.002870	2.132375	12.739749	1698.486030	349.722415	963.895318
min	-124.350000	32.540000	1.000000	2.000000	2.000000	1.000000
25%	-121.740000	33.950000	19.000000	1436.000000	301.000000	806.000000
50%	-118.460000	34.240000	30.000000	2079.500000	437.000000	1166.000000
75%	-118.060000	37.720000	39.000000	3042.000000	643.000000	1742.000000
max	-114.310000	41.950000	52.000000	14125.000000	2823.000000	7603.000000

In []:

The code below creates a new column with the distance a block is from a major city

In [241]:

Out[241]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	p
count	20258.000000	20258.000000	20258.000000	20258.000000	20258.000000	20258.000000
mean	-119.575046	35.636977	28.754418	2537.309409	520.139599	137.000000
std	2.003614	2.137204	12.523427	1750.309365	351.709278	92.000000
min	-124.350000	32.540000	1.000000	2.000000	2.000000	1.000000
25%	-121.800000	33.930000	18.000000	1448.000000	295.000000	78.000000
50%	-118.500000	34.260000	29.000000	2119.000000	433.000000	116.000000
75%	-118.010000	37.720000	37.000000	3114.000000	641.000000	176.000000
max	-114.310000	41.950000	52.000000	14125.000000	2823.000000	765.000000

The cell below creates some new dataframes containing the data in the area of los angeles and san francisco

In [221]:

```
longlostangel = housing[ (housing['longitude'] <= -117) & (housing['longitude'] >= -119)
lostangel = longlostangel[ (longlostangel['latitude'] >= 33) & (longlostangel['latitude'] <= 37)]

santa = housing[ (housing['longitude'] <= -121) & (housing['longitude'] >= -123)]
san = santa[ (santa['latitude'] >= 37) & (santa['latitude'] <= 39)]
san.head(10)
```

Out[221]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	household
0	-122.23	37.88	41	880	129.0	322	1
1	-122.22	37.86	21	7099	1106.0	2401	1
2	-122.24	37.85	52	1467	190.0	496	1
3	-122.25	37.85	52	1274	235.0	558	1
4	-122.25	37.85	52	1627	280.0	565	1
5	-122.25	37.85	52	919	213.0	413	1
6	-122.25	37.84	52	2535	489.0	1094	1
7	-122.25	37.84	52	3104	687.0	1157	1
8	-122.26	37.84	42	2555	665.0	1206	1
9	-122.25	37.84	52	3549	707.0	1551	1

Bedlessness


Import shapefile

In [234]:

```
from shapely.geometry import Point, Polygon
import os
import seaborn as sns
import shapely.ops
import netCDF4
from pyproj import Transformer
import matplotlib.patches as mpatches
import matplotlib.lines as mlines
import geopandas as gpd
```

In [235]:

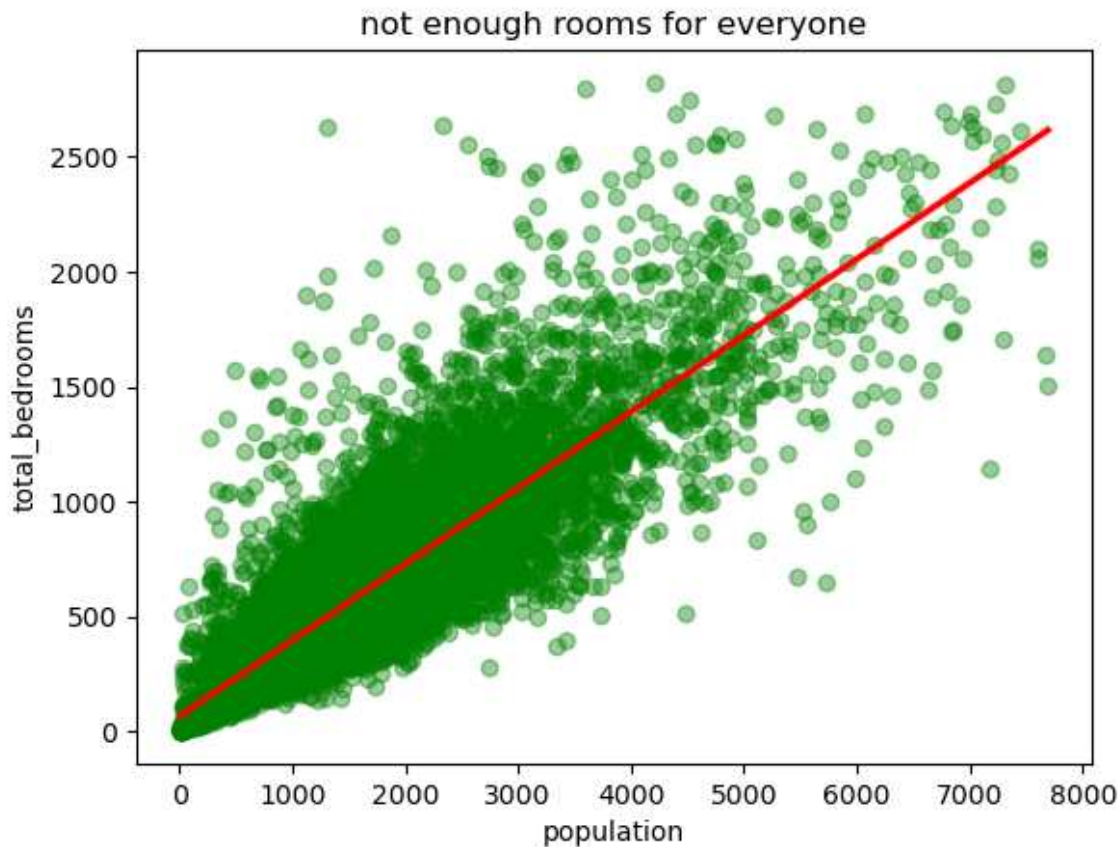
```
os.environ['SHAPE_RESTORE_SHX'] = 'YES' #for some reason I got an error that told me to
#importing our stuff
county_shapefile = gpd.read_file('CA_Counties_TIGER2016.shp') #shape file for Cali map
county_shapefile['geometry'] = county_shapefile['geometry'].apply(lambda geom: shapely.o
```



Regression Plot of population to bedrooms

In [236]:

```
sns.regplot(data=housing, x=housing['population'], y=housing['total_bedrooms'], line_kws=
plt.title('not enough rooms for everyone')
plt.show()
p1 = np.poly1d(np.polyfit(housing['population'], housing['total_bedrooms'], 1))
pv = p1(housing['population'])
r2 = r2_score(housing['total_bedrooms'], pv)
print(p1)
print(r2)
```



```
0.3317 x + 64.14
0.7536994142357631
```

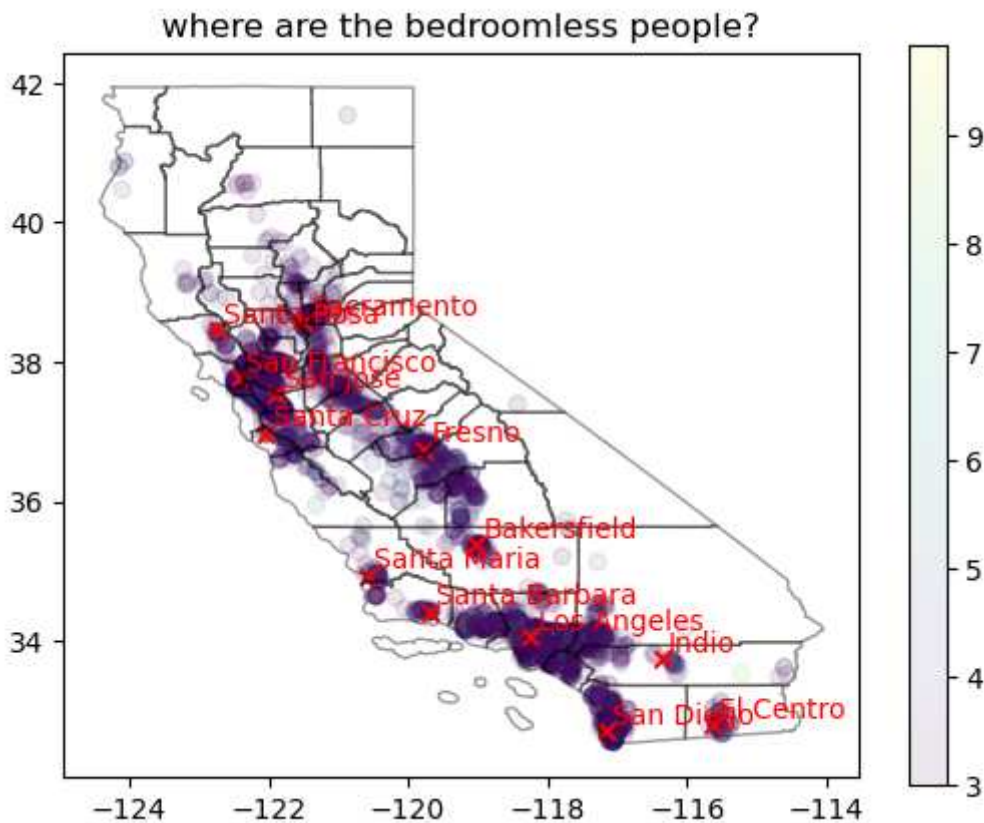
On average 70% of Californians either share a bedroom or don't have a bedroom to sleep in. A significant amount of these bedless people reside in the Los Angeles area and the San Francisco area as seen in the map below showing the locations where the population and total number of bedrooms exceeds 4:1. We can also see from the regression plot showing the relationship between distance from major city and homelessness, that bedlessness is more frequent the closer you get to a city

In [237]:

```

calirooms = housing[(housing['population'] >= 3*housing['total_bedrooms'])]
#make sure there are no outliers so that we can visualise better:
upper_fence_rooms_total_bedrooms = np.percentile(calirooms['total_bedrooms'], 75) + 1.5
upper_fence_rooms_population = np.percentile(calirooms['population'], 75) + 1.5 * (np.pe
lower_fence_rooms_total_bedrooms = np.percentile(calirooms['total_bedrooms'], 25) - 1.5
lower_fence_rooms_population = np.percentile(calirooms['population'], 25) - 1.5 * (np.pe
calirooms = calirooms[(calirooms['population'] >= lower_fence_rooms_population) & (calirooms['total_bedrooms'] >= lower_fence_rooms_total_bedrooms)]
#now lets plot it with respect to a map
fig, ax = plt.subplots()
plt.scatter(x=calirooms['longitude'],y=calirooms['latitude'],alpha=0.1,c=calirooms['population'])
plt.colorbar()
plt.scatter(x=california_cities['longitude'],y=california_cities['latitude'],alpha = 1,marker='x')
for i in range(california_cities.shape[0]):
    plt.text(x=california_cities.longitude[i]+0.1,y=california_cities.latitude[i]+0.1,s=california_cities.city[i])
county_shapefile.plot(ax=ax, color='None', edgecolor='black',alpha=0.4)
plt.title('where are the bedroomless people?')
plt.show();

```




In [238]:

```


calirooms = san[(san['population'] >= 3*san['total_bedrooms'])]
#make sure there are no outliers so that we can visualise better:
#upper_fence_rooms_total_bedrooms = np.percentile(calirooms['total_bedrooms'], 75) + 1.5
#upper_fence_rooms_population = np.percentile(calirooms['population'], 75) + 1.5 * (np.p
#lower_fence_rooms_total_bedrooms = np.percentile(calirooms['total_bedrooms'], 25) - 1.5
#lower_fence_rooms_population = np.percentile(calirooms['population'], 25) - 1.5 * (np.p
#calirooms = calirooms[(calirooms['population'] >= lower_fence_rooms_population) &(calir
#now lets plot it with respect to a map
fig, ax = plt.subplots()
plt.scatter(x=calirooms['longitude'],y=calirooms['latitude'],alpha=0.5,c=calirooms['popu
plt.colorbar()
plt.title('San Francisco & Santa Cruz area')
plt.plot(-122.030797,37.090017, marker="o", markersize=25, markeredgecolor="red", marker
plt.text(-122.230797,37.154117, 'Santa Cruz',c='red')
plt.plot(-121.889096, 37.354960, marker="o", markersize=50, markeredgecolor="red", marke
plt.text(-122.059096, 37.354960, 'San Jose',c='red')
plt.plot(-122.469417,37.704931, marker="o", markersize=40, markeredgecolor="red", marker
plt.text(-122.306034,37.883974, 'San Francisco',c='red')
plt.show();

```



In [239]:


```
calirooms = lostangel[(lostangel['population'] >= 3*lostangel['total_bedrooms'])]
#make sure there are no outliers so that we can visualise better:
upper_fence_rooms_total_bedrooms = np.percentile(calirooms['total_bedrooms'], 75) + 1.5
upper_fence_rooms_population = np.percentile(calirooms['population'], 75) + 1.5 * (np.pe
lower_fence_rooms_total_bedrooms = np.percentile(calirooms['total_bedrooms'], 25) - 1.5
lower_fence_rooms_population = np.percentile(calirooms['population'], 25) - 1.5 * (np.pe
calirooms = calirooms[(calirooms['population'] >= lower_fence_rooms_population) &(calirooms['total_bedrooms'] >= lower_fence_rooms_total_bedrooms)]
#now lets plot it with respect to a map
fig, ax = plt.subplots()
plt.scatter(x=calirooms['longitude'],y=calirooms['latitude'],alpha=0.5,c=calirooms['population'])
plt.colorbar()
#plt.legend()
plt.title('Los Angeles area')
plt.plot(-118.243686,34.052233, marker="o", markersize=70, markeredgecolor="red", markerfacecolor="yellow")
plt.text(-118.455,34.022233, 'Los Angeles',c='red')
plt.show();
```



In [249]:

```
plt.title('bedlessness in relation to proximity to city')
p2 = np.poly1d(np.polyfit(housing['distance_to_closest_city'], housing['bedless'], 2))
housefit = np.linspace(housing['distance_to_closest_city'].min(), housing['distance_to_c
plt.scatter(x=housing['distance_to_closest_city'], y=housing['bedless'], alpha=0.5, c=hous
plt.plot(housefit, p2(housefit), label="Average", c='red')
pv = p2(housing['distance_to_closest_city'])
r2 = r2_score(housing['bedless'], pv)
print(r2)
plt.show()
```

0.014904420412508168



Median income and house value as approaching city

Median Income

In [48]:

```

caliincomebin = pd.cut(housing['median_income'], 4, precision=2)
calivalubin = pd.cut(housing['median_house_value'], 4, precision=2)

conditions = [
    (housing['median_income'] <= 2.24),
    (housing['median_income'] > 2.24) & (housing['median_income'] <= 4.48),
    (housing['median_income'] > 4.48) & (housing['median_income'] <= 8),
    (housing['median_income'] > 8)
]
fence_house_value = np.percentile(housing['median_house_value'], 75) + 1.5 * (np.percent
Vconditions = [
    (housing['median_house_value'] <= np.percentile(housing['median_house_value'], 25)),
    (housing['median_house_value'] > np.percentile(housing['median_house_value'], 25)) &
    (housing['median_house_value'] > np.percentile(housing['median_house_value'], 75)) &
    (housing['median_house_value'] > fence_house_value)
]

status = ['poor', 'middle_class', 'wealthy', 'uber_rich']
housing['status'] = np.select(conditions, status)
house_bins = ['Low_Value', 'Medium_Value', 'High_Value', 'Highest_Value']
housing['house_bin'] = np.select(Vconditions, house_bins)
housing.head()

```

C:\Users\benjo\AppData\Local\Temp\ipykernel_4648\19209354.py:19: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
housing['status'] = np.select(conditions, status)
```

C:\Users\benjo\AppData\Local\Temp\ipykernel_4648\19209354.py:21: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
housing['house_bin'] = np.select(Vconditions, house_bins)
```

Out[48]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	househ
0	-122.23	37.88	41	880	129.0	322	
1	-122.22	37.86	21	7099	1106.0	2401	
2	-122.24	37.85	52	1467	190.0	496	
3	-122.25	37.85	52	1274	235.0	558	
4	-122.25	37.85	52	1627	280.0	565	

In [47]:

```

fig, ax = plt.subplots()

status_colors = {'poor': 'orange', 'middle_class': 'green', 'wealthy': 'blue', 'uber_ric

# Map the categories to colors for each data point
statuscolors = housing['status'].map(status_colors)

#county_shapefile.plot(ax=ax, color='white', edgecolor='black')

plt.scatter(x=housing['longitude'],y=housing['latitude'],alpha = (35+((1.066**((100*(hous

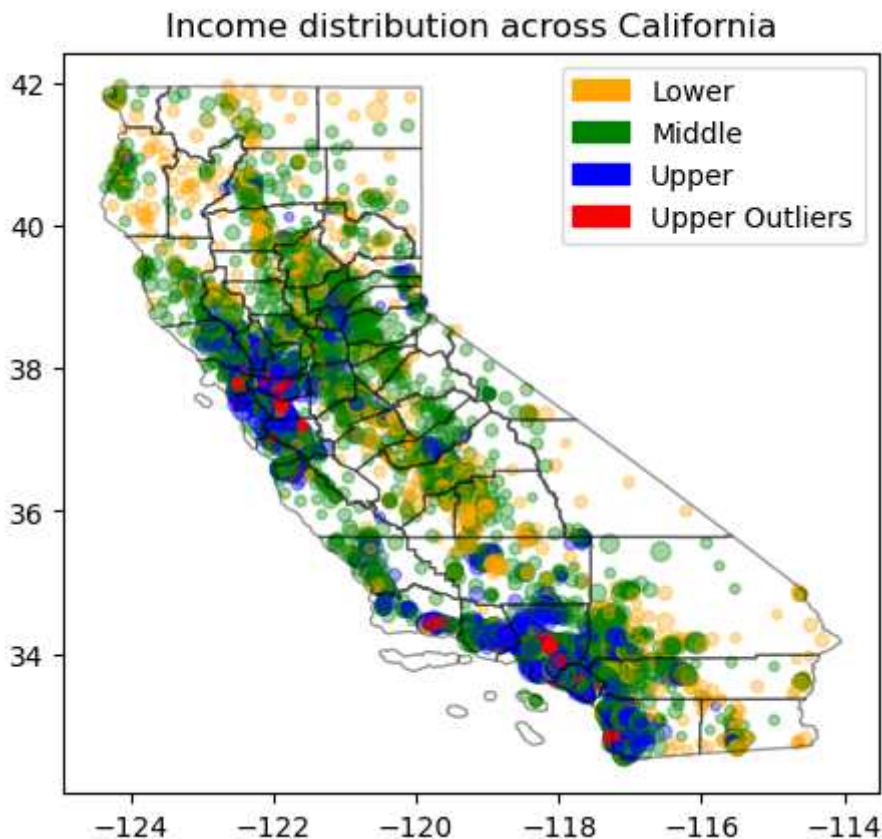
#plt.scatter(x=california_cities['longitude'],y=california_cities['latitude'],alpha = 1,
#for i in range(california_cities.shape[0]):
#    plt.text(x=california_cities.longitude[i]+0.1,y=california_cities.latitude[i]+0.1,s
county_shapefile.plot(ax=ax, color='None', edgecolor='black',alpha=0.4)

plt.title("Income distribution across California")

outliers = mpatches.Patch(color='red', label='Upper Outliers')
poor = mpatches.Patch(color='orange', label='Lower')
mid = mpatches.Patch(color='green', label='Middle')
upper = mpatches.Patch(color='blue', label='Upper')
plt.legend(handles=[poor,mid,upper,outliers])

plt.show();


```



In [196]:

```
plt.title('Median Income in relation to proximity to city')
p2 = np.poly1d(np.polyfit(housing['distance_to_closest_city'], housing['median_income'],
housefit = np.linspace(housing['distance_to_closest_city'].min(), housing['distance_to_c
plt.scatter(x=4-housing['distance_to_closest_city'], y=housing['median_income'], c='gree
plt.plot(4-housefit, p2(housefit), label="Model", c='red')
print(r2)
plt.legend()
plt.xlabel('Distance to Closest City (Inverted)')
plt.ylabel('Median Income')
plt.show()
```

0.09504330310432407



In []:

Median House price

In [49]:

```

fig, ax = plt.subplots()

house_bin_colors = {'Low_Value': 'orange', 'Medium_Value': 'green', 'High_Value': 'blue'}

# Map the categories to colors for each data point
house_bin_colors = housing['house_bin'].map(house_bin_colors)

plt.scatter(x=housing['longitude'],y=housing['latitude'],alpha = (35+((1.066**((100*(housing['house_value']-housing['house_value'].min())/(housing['house_value'].max()-housing['house_value'].min()))**2))))

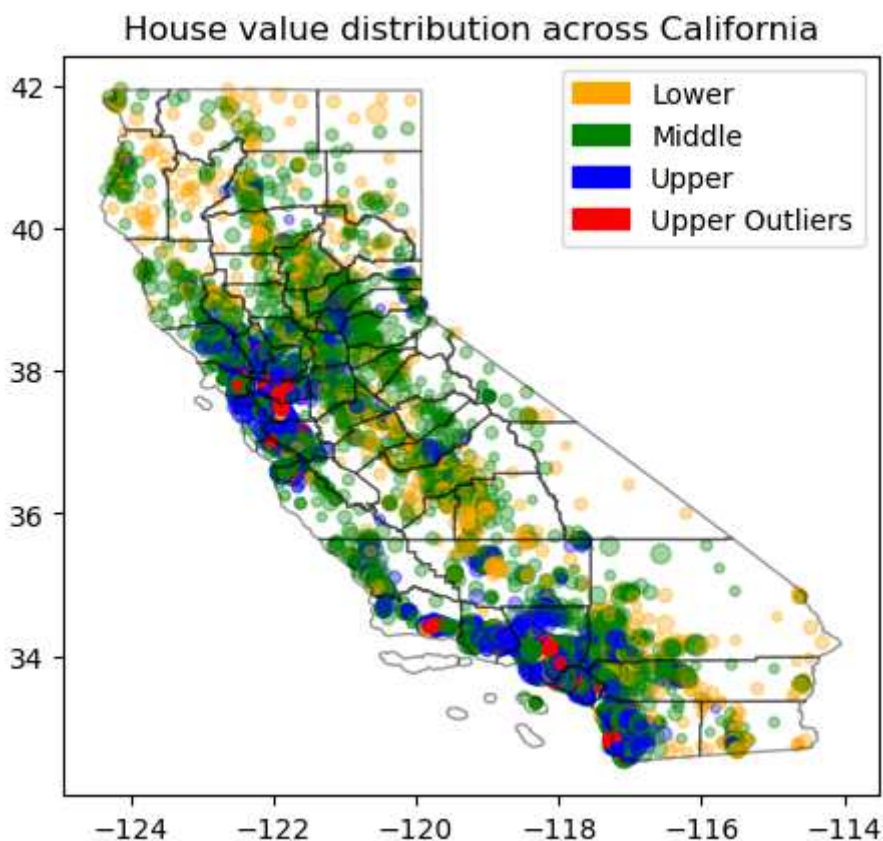
county_shapefile.plot(ax=ax, color='None', edgecolor='black',alpha=0.4)

plt.title("House value distribution across California")

Highest_Value = mpatches.Patch(color='red', label='Highest_Value')
Low_Value = mpatches.Patch(color='orange', label='Low_Value')
Medium_Value = mpatches.Patch(color='green', label='Medium_Value')
High_Value = mpatches.Patch(color='blue', label='High_Value')
plt.legend(handles=[poor,mid,upper,outliers])

plt.show();

```




In []:

In [215]:

```
plt.title('Median House value in relation to proximity to city')
p2 = np.poly1d(np.polyfit(housing['distance_to_closest_city'], housing['median_house_val
housefit = np.linspace(housing['distance_to_closest_city'].min(), housing['distance_to_c
plt.scatter(x=4-housing['distance_to_closest_city'], y=housing['median_house_value'],c='
plt.plot(4-housefit,p2(housefit),label="Average",c='red')
pv = p2(housing['distance_to_closest_city'])
r2 = r2_score(housing['median_house_value'],pv)
print(r2)
plt.show()
```

0.09123367804765248



In []: